

CA1 - PartA

# Convolutional Neural Network

Problem Statement:

Predicting Vegetables

May 20 2025

Kaung Myat San, P2408655

# Agenda

- Background Research & Data Analysis
- Data Augmentation
- Baseline Model
- Hyper-Tuning
- Improving Hyper-Tuned





Background  
Research & Data  
Analysis



# Dataset Overview

- Image Size: 224 x 224 pixels (3 color channels- rgb)
- Number of class: 11 classes (13 vegetables)
- Classes:
  - Bean
  - Bitter Gourd
  - Bottle Gourd and Cucumber
  - Brinjal
  - Broccoli and Cauliflower
  - Cabbage
  - Capsicum
  - Carrot and Radish
  - Potato
  - Pumpkin
  - Tomato



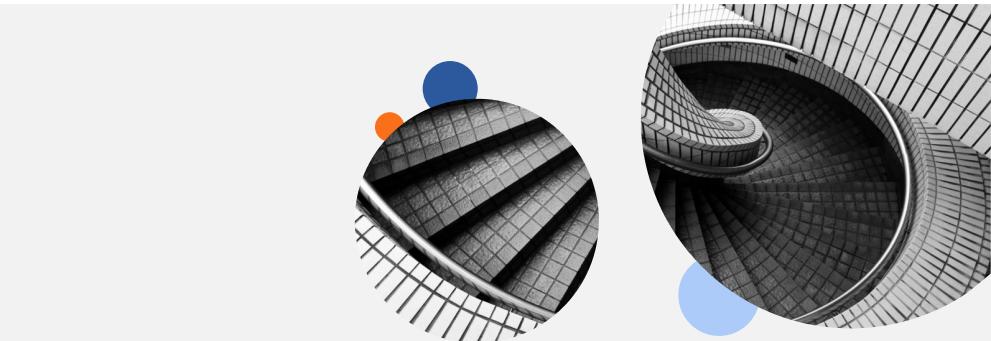
# Data Analysis

## Irregularities

- Test Folder:
  - Pumpkin and Tomato Folders are swapped
- Train Folder:
  - Bean Folder: 0001-0020, 0049, 0050.jpg are carrots

## How to fix

- Assigned it as the correct label when processing and loading image
- Ensure correct label for each vegetable folder



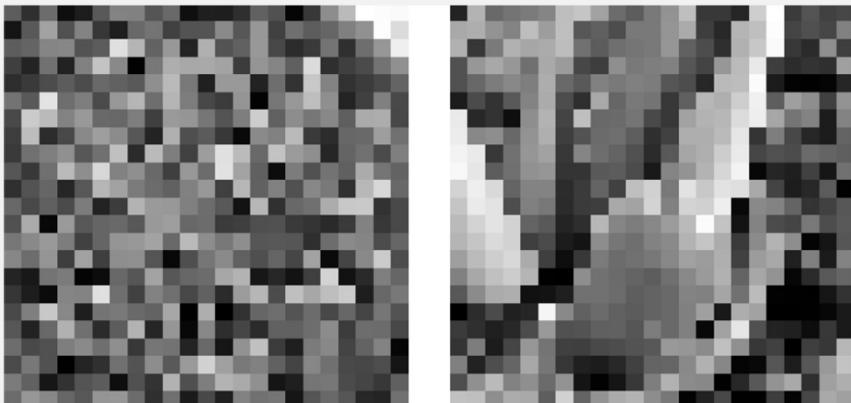
## Metrics to Evaluate

- **Loss** – metric to monitor during training
- **Recall** – Reduce false negatives (catch all actual vegetables)
- **Accuracy** – Overall Correct Predictions
- **Precision** – Reduce false positives (avoid wrong vegetable predictions)
- **Focus**: Loss, Recall & Accuracy
- **Reason**: shows a complete view of the model's performance

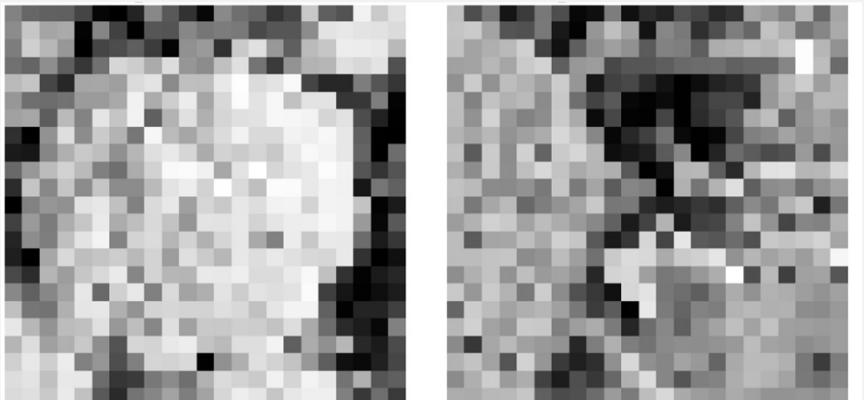
# Why Group Some Vegetables?



Bottle Gourd & Cucumber



Broccoli & Cauliflower



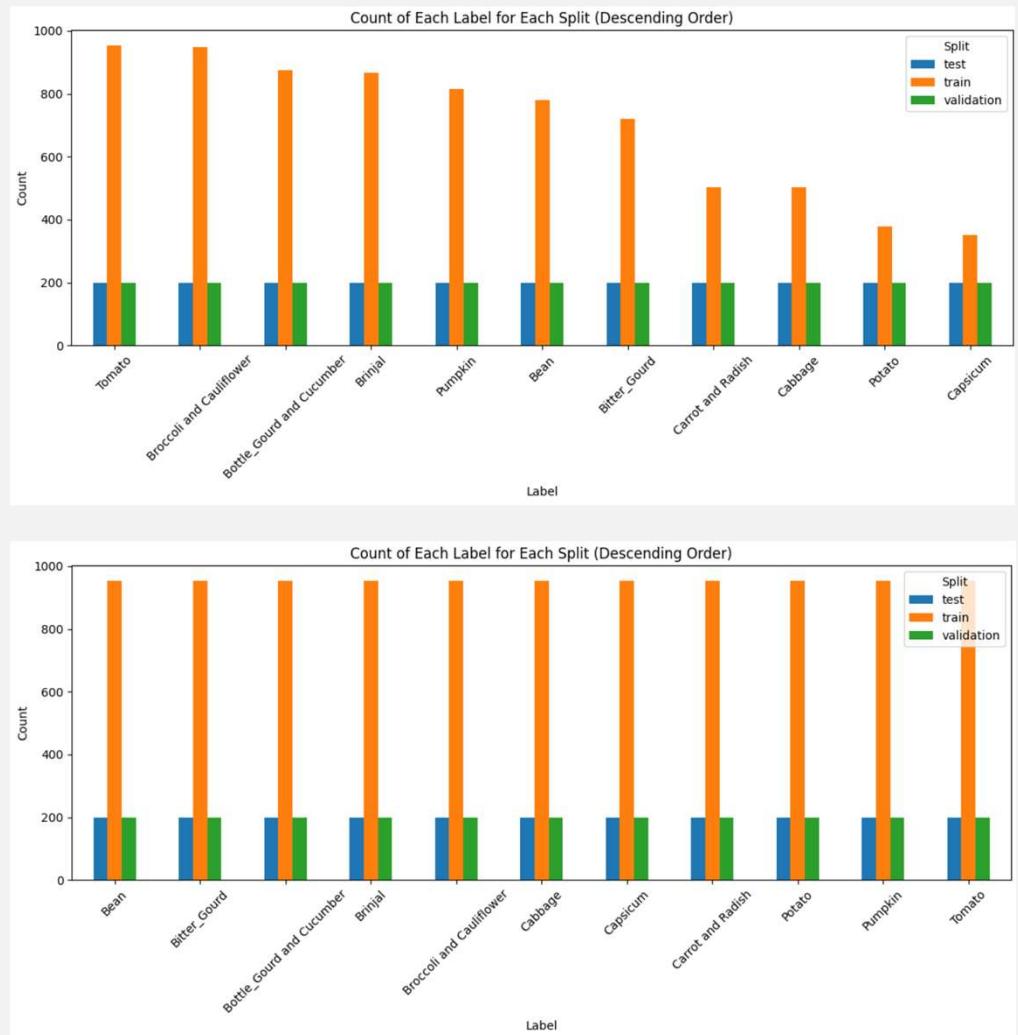
- Can you differentiate between these 2 vegetables?
- These 2 vegetables looked very similar with gray-scaled and challenging for our model
- Therefore, these vegetables will be grouped together

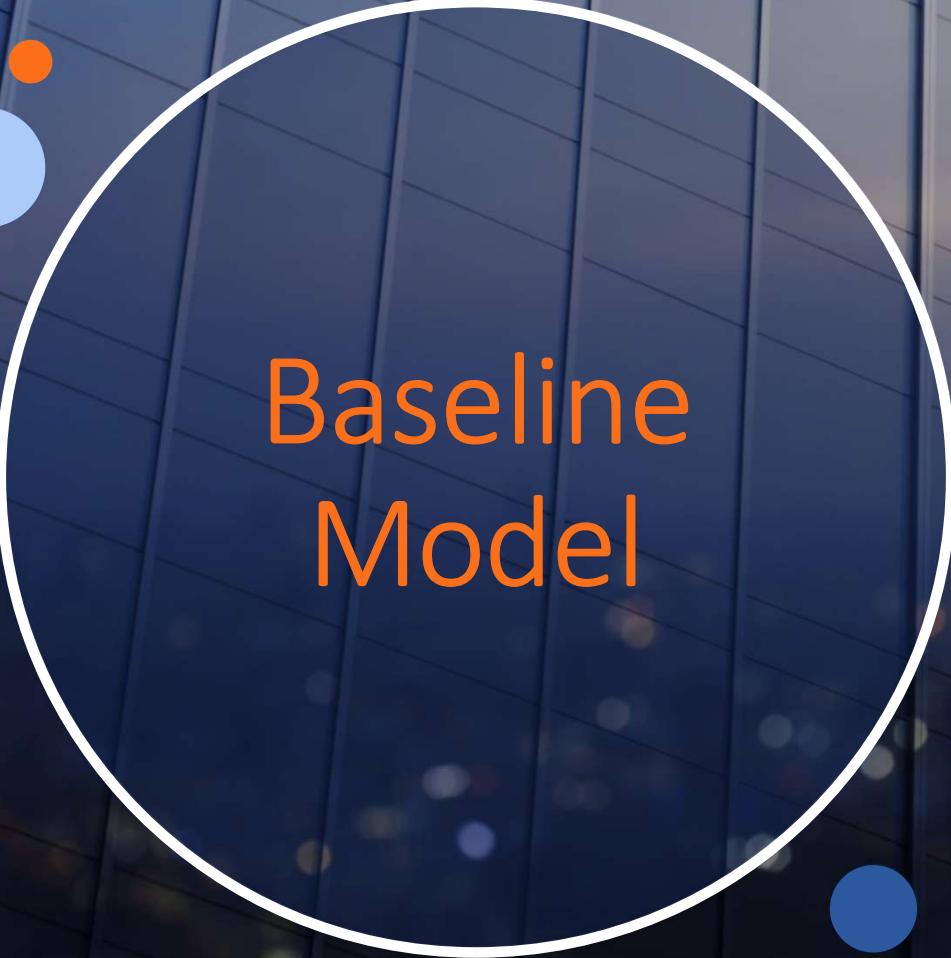


Data  
Processing

# Data Processing

- Convert to gray-scale and corresponding size
- Convert to array
- Normalized the array
- Convert into one-hot encoding
- Removed duplicates – 17
- Perform Data Transformation to create a more balanced dataset
- Generate Images so all classes have same amount as the max.





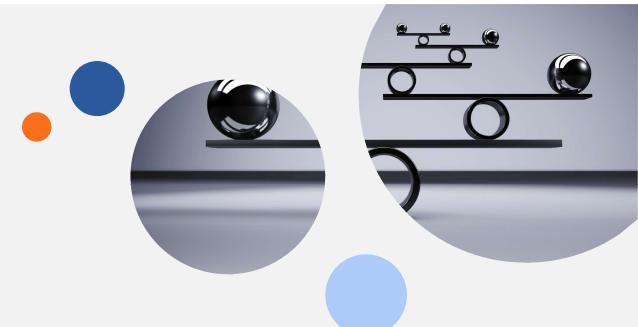
Baseline  
Model

# Base Model Layers

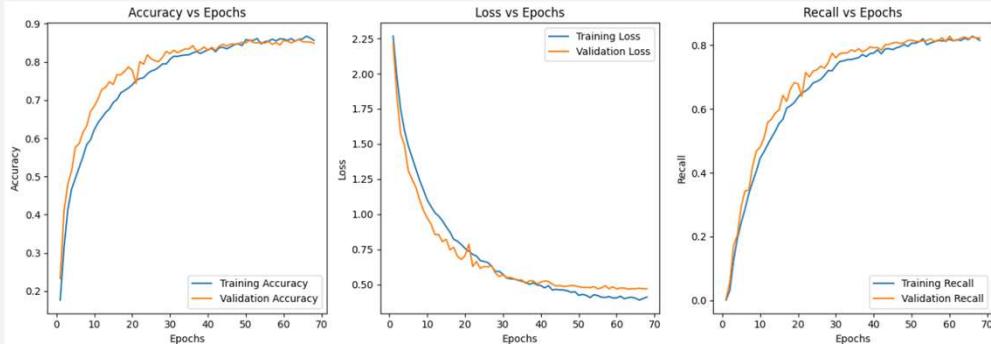
- Compile
  - optimizer = Adam(learning\_rate = 0.001)
  - loss = "categorical\_crossentropy"
- Fitting
  - Early Stopping - monitor= val\_loss, patience = 5, restore\_best\_weights = true
  - reduce\_lr – monitor = val\_loss, factor = 0.5, patience = 3, min\_lr = 1e-6
  - epochs = 100
  - batch\_size = 32

Layers	Parameters
Conv2D	Filters = 32
MaxPooling2D	pool_size = (2,2)
Dropout	0.25
Conv2D	Filters = 64
MaxPooling2D	pool_size = (2,2)
Dropout	0.25
Flatten	0.25
Dense	128, 'relu'
Dropout	0.5
Dense	11, 'softmax'

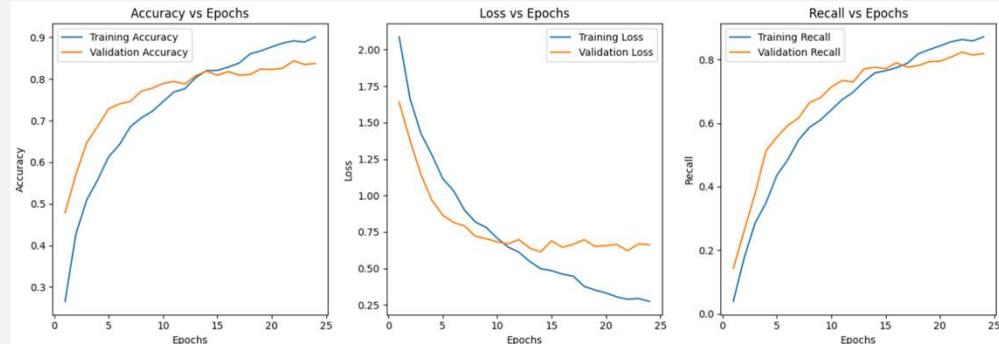
# Balanced Vs Imbalanced (101x101)



## Balanced

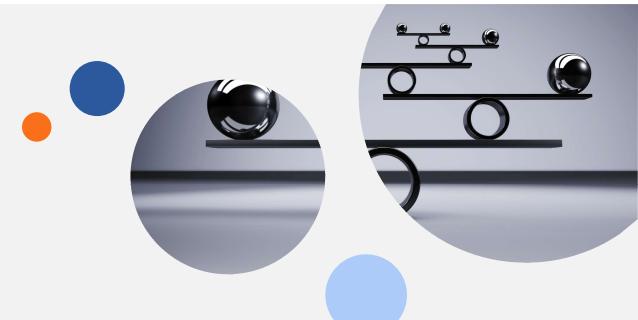


## Imbalanced

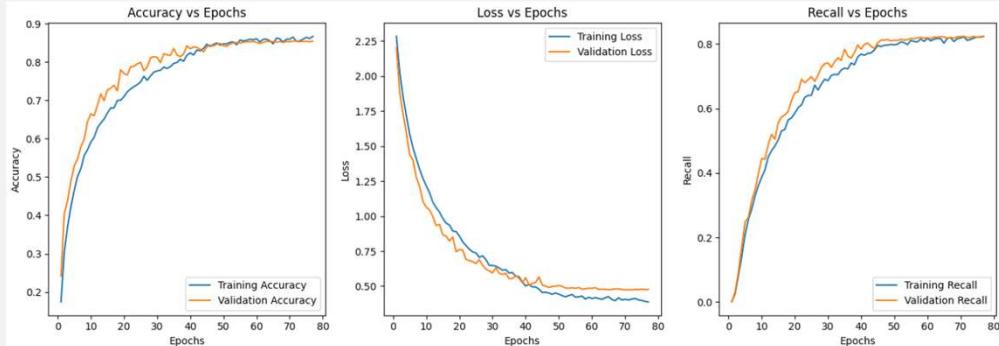


- Precision: 86%
  - Recall: 85%
  - Accuracy: 85.23%
- Precision: 82%
  - Recall: 82%
  - Accuracy: 82.23%

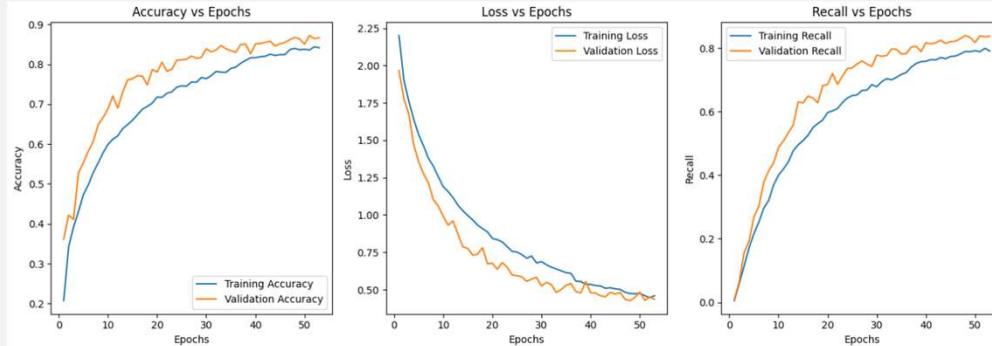
# Balanced Vs Imbalanced (23x23)



## Balanced

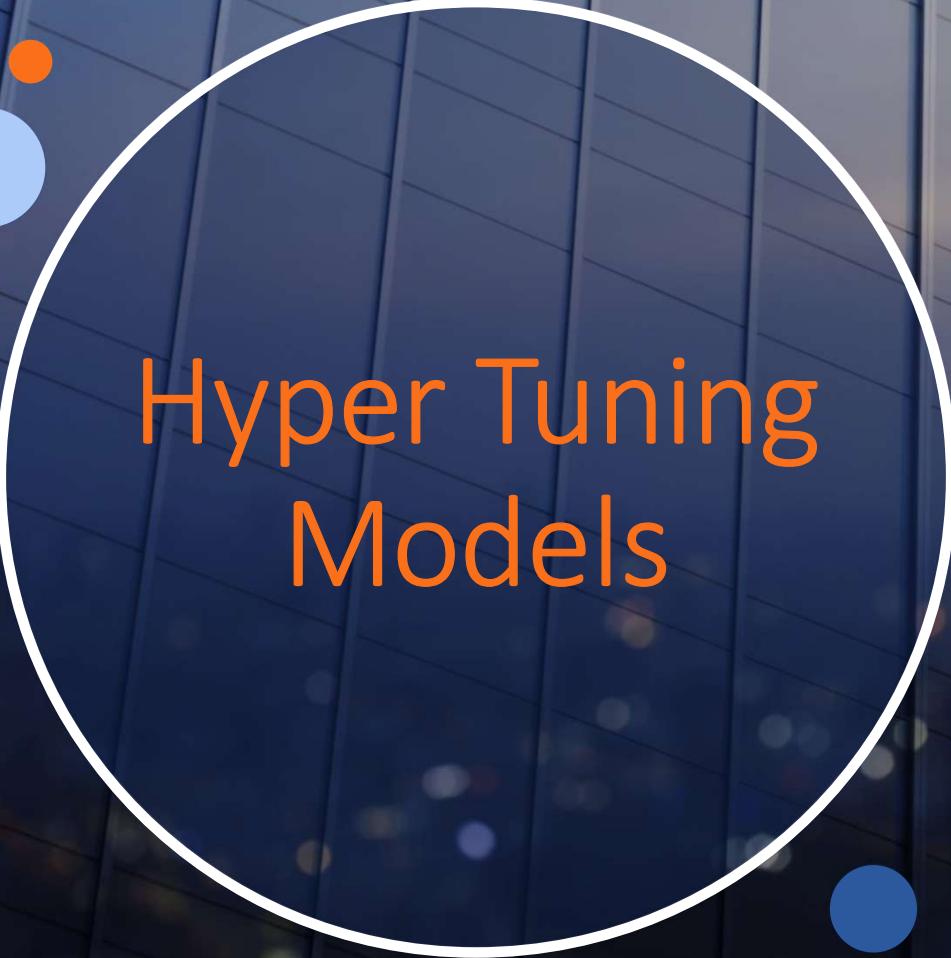


## Imbalanced



- Precision: 87%
- Recall: 86%
- Accuracy: 86.27%

- Precision: 87%
- Recall: 86%
- Accuracy: 86.41%



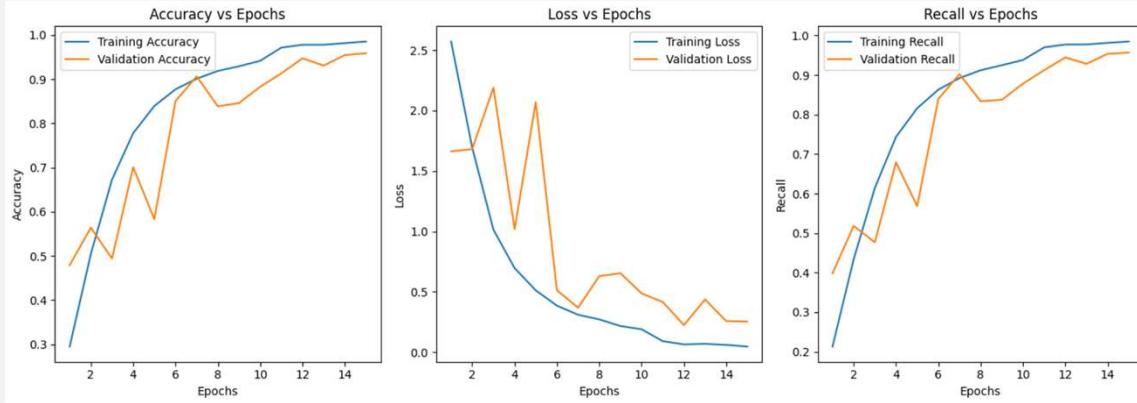
Hyper Tuning  
Models



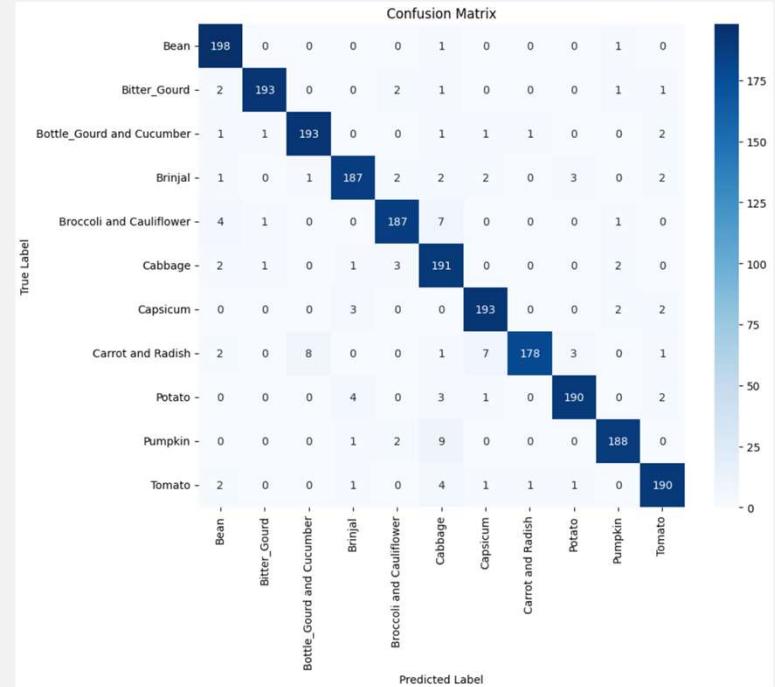
# Hyper Tuning Parameters

Parameters	Details	Parameters	Details
Conv_filters	The order of dimension of filters in Conv2D	final_activation	Activation function for the output layer
kernel_size	The order of list of convolutional windows	optimizer	Optimizer to use for compiling the model
pool_size	Pool size for MaxPooling2D layers	batch_norm	Whether to include BatchNormalization layers
dense_units	Number of units in the dense layer	learning_rate	Learning rate for the optimizer
dropout_rates	List of dropout rates for each Dropout layer	no_of_conv_layers	Number of Conv2D layers to stack for each filter size
activation	Activation function for Conv2D and Dense layers	l2_regularization	L2 regularization factor for Conv2D layers

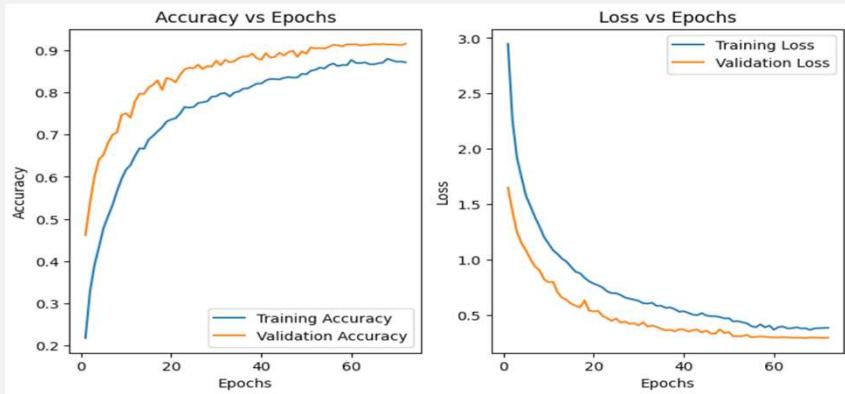
# 101x101 Best Hyper Tuned Model



- Test Recall: 0.95
  - Test Precision: 0.95
  - Test Accuracy: 0.95
- Area for Improvement
- Over-Fitting:** Reduce dropout slightly
  - Misclassification:** Augment more data for Poorly Classified Class



# 23x23 Best Hyper Tuned Model



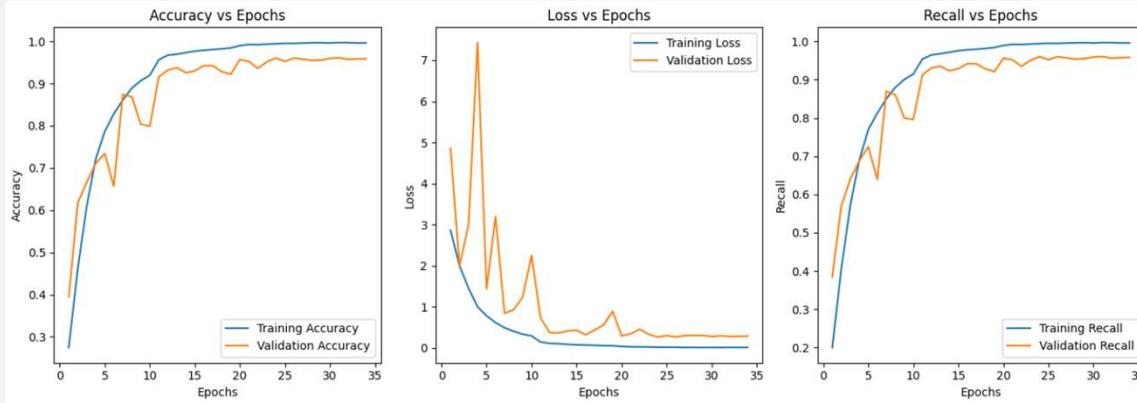
- Test Recall: 0.90
- Test Precision: 0.90
- Test Accuracy: 0.90

## Area for Improvement

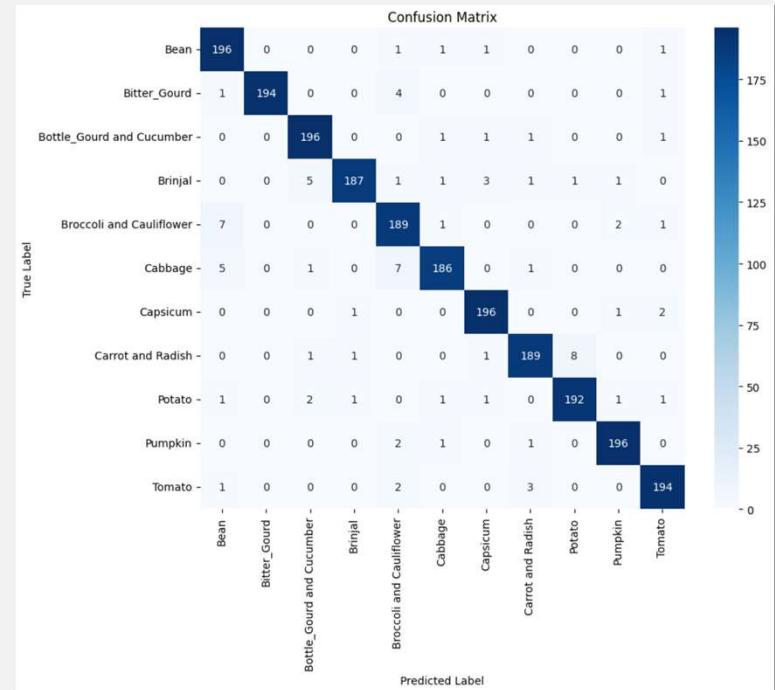
- Misclassification:** Augment more data for Poorly Classified Class
  - Broccoli & Cauliflower
  - Cabbage
  - Tomato
  - Potato
  - Bitter Gourd & Cucumber
  - Bean

		Confusion Matrix											
		True Label	Bean	Bitter_Gourd	Bottle_Gourd and Cucumber	Brinjal	Broccoli and Cauliflower	Cabbage	Capsicum	Carrot and Radish	Potato	Pumpkin	Tomato
Predicted Label	Bean	187	1	3	0	3	1	0	1	0	1	3	
	Bitter_Gourd	9	176	0	0	8	1	0	0	0	1	5	
	Bottle_Gourd and Cucumber	1	1	191	2	1	1	0	0	0	1	0	
	Brinjal	0	3	2	188	0	1	1	2	1	2	0	
	Broccoli and Cauliflower	13	4	4	1	164	3	1	0	0	3	7	
	Cabbage	2	2	0	1	18	171	1	0	0	4	1	
	Capsicum	1	1	2	1	0	0	191	1	0	2	1	
	Carrot and Radish	1	0	1	1	3	0	0	179	11	0	4	
	Potato	2	0	2	7	0	3	3	4	175	1	3	
	Pumpkin	1	0	0	3	4	1	1	1	0	186	3	
	Tomato	2	2	2	2	6	1	1	8	1	2	173	

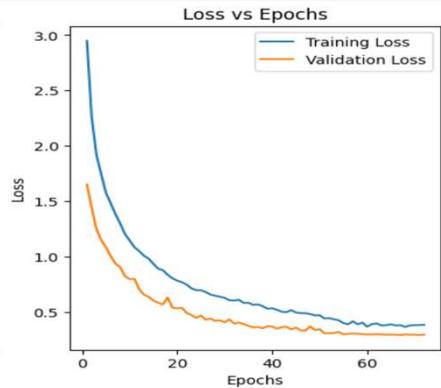
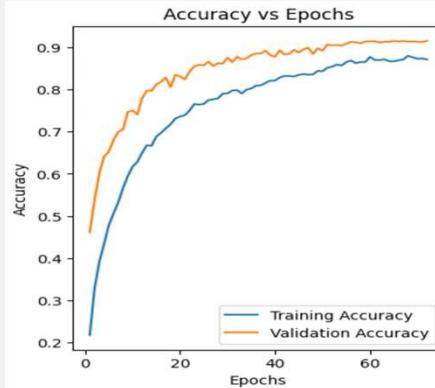
# 101x101 Final Best Model



- Test Recall: 0.96
- Test Precision: 0.96
- Test Accuracy: 0.96

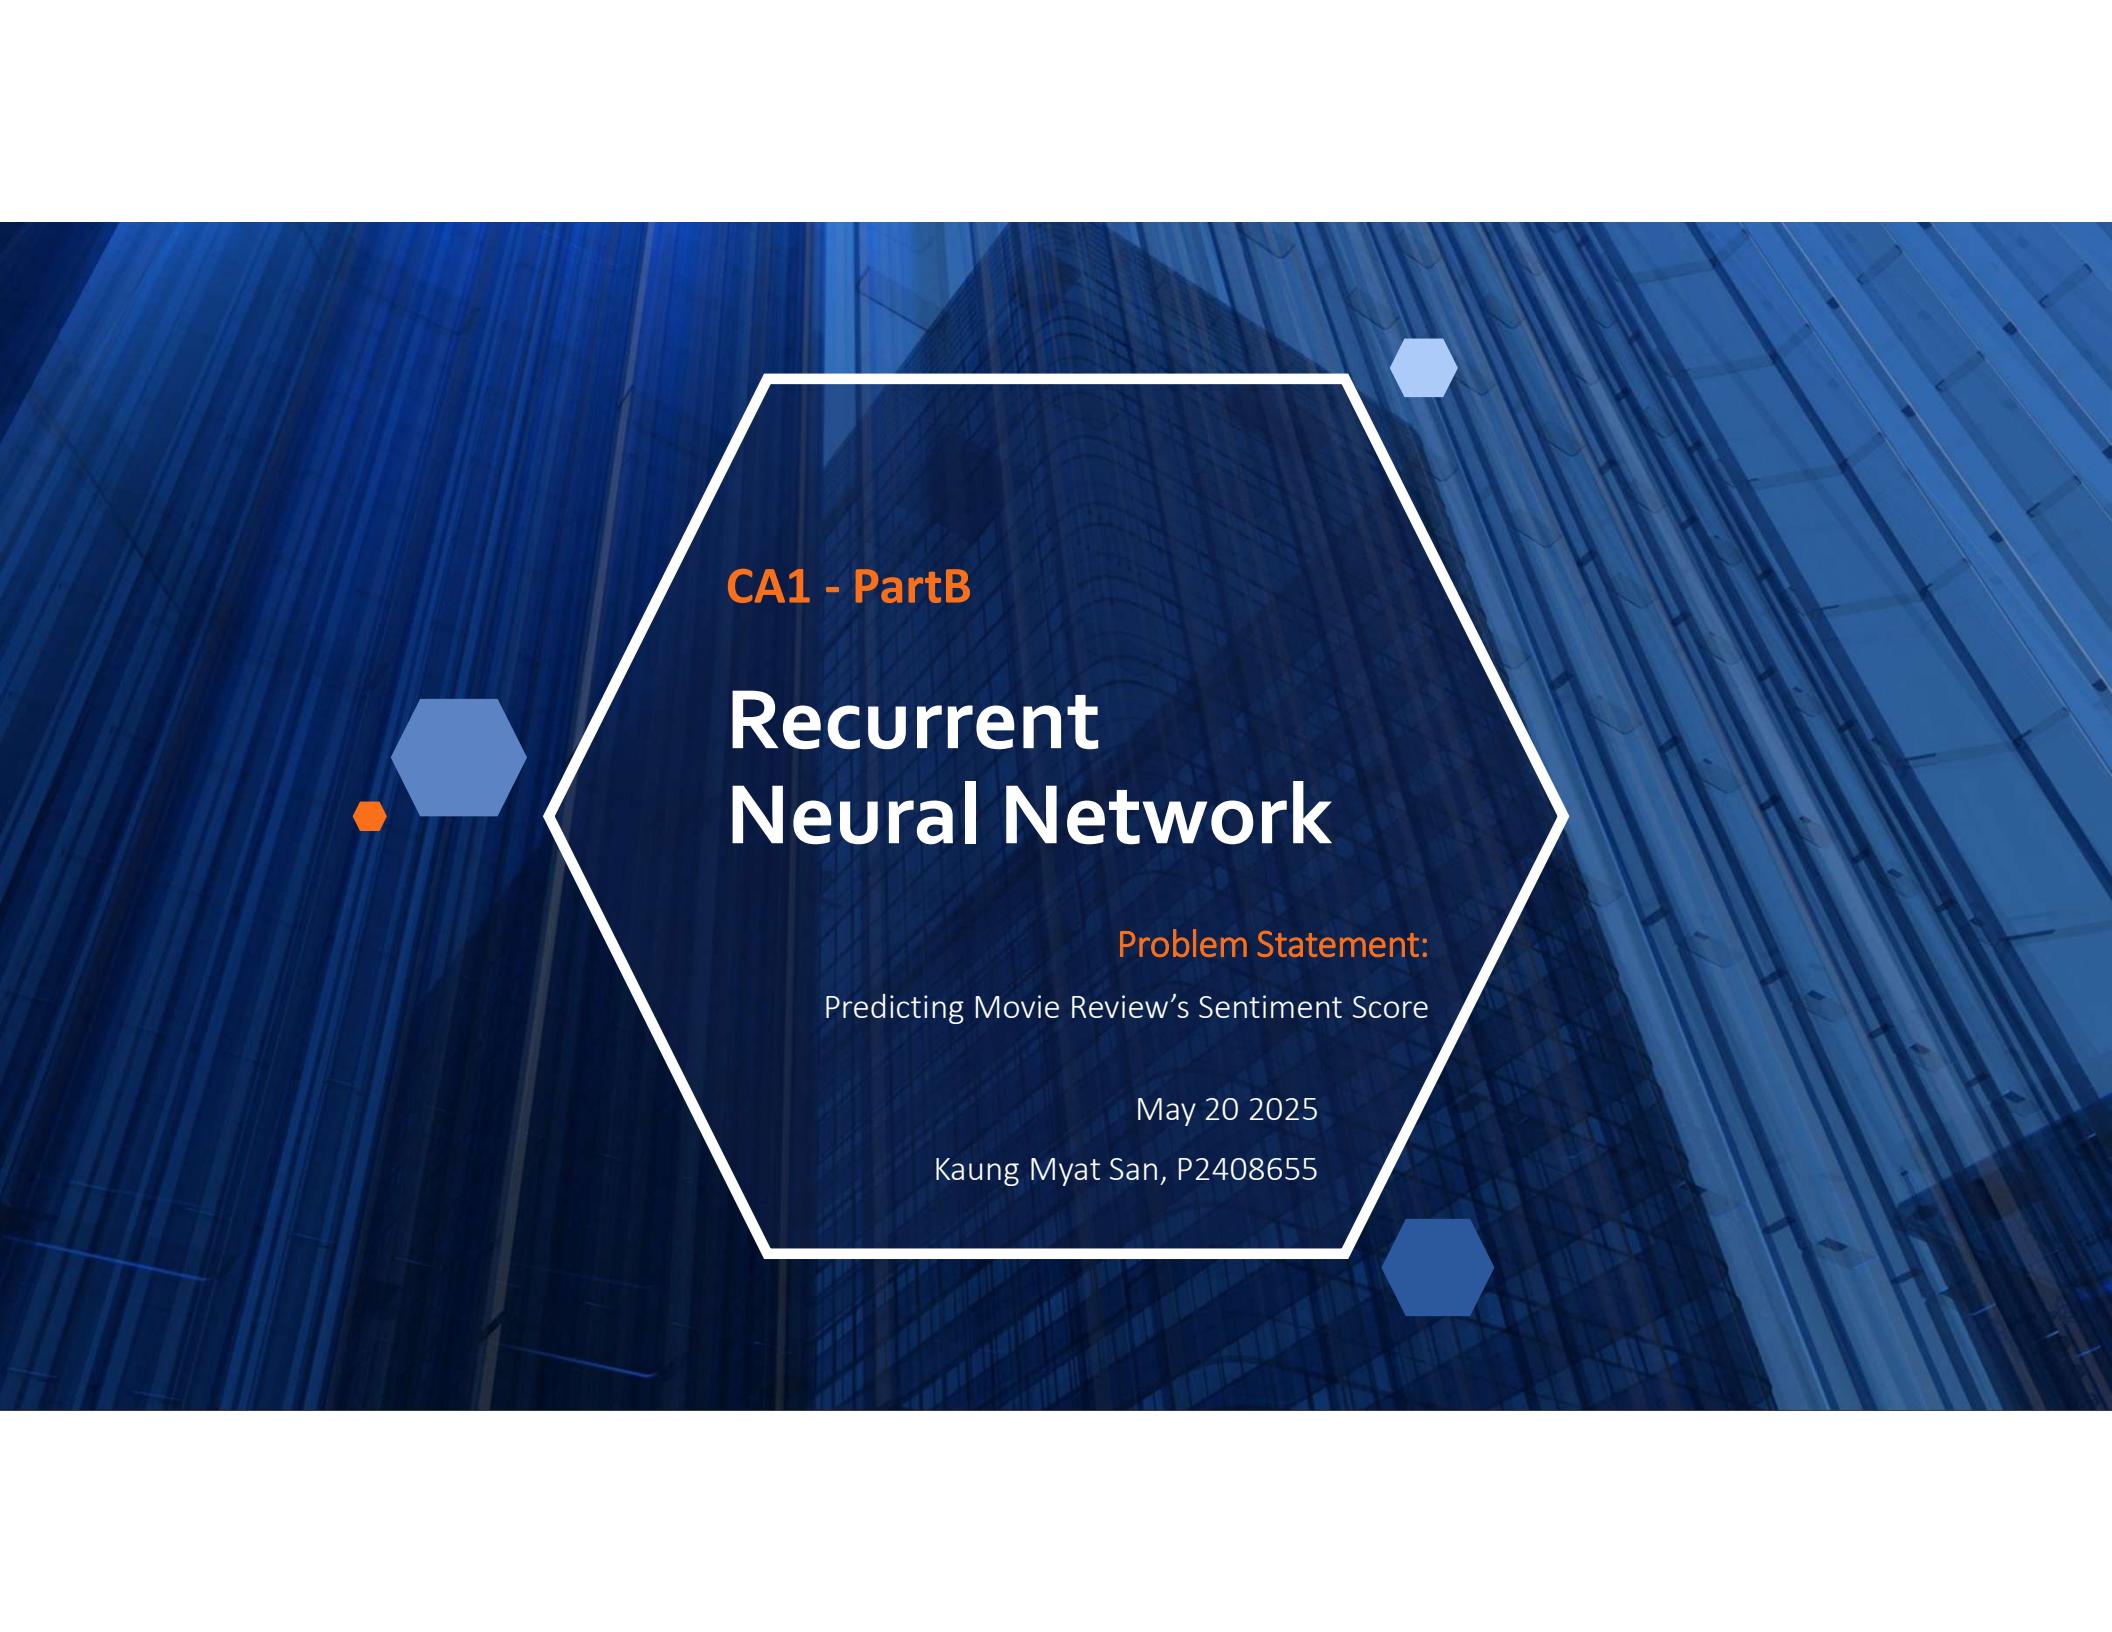


# 23x23 Final Best Model



- Didn't Improve after trying Data Augmentation and Data Removal based on Confusion Matrix
- Reason: Augmenting a 23 by 23 images might cause them to lose distinctive features, making different vegetables similar
- Test Recall: 0.90
- Test Precision: 0.90
- Test Accuracy: 0.90

Confusion Matrix												
True Label	Bean	Bitter_Gourd	Bottle_Gourd and Cucumber	Brinjal	Broccoli and Cauliflower	Cabbage	Capsicum	Carrot and Radish	Potato	Pumpkin	Tomato	
	187	1	3	0	3	1	0	1	0	1	3	175
	9	176	0	0	8	1	0	0	0	1	5	150
	1	1	191	2	1	1	0	0	1	0	2	125
	0	3	2	188	0	1	1	2	1	2	0	100
	13	4	4	1	164	3	1	0	0	3	7	75
	2	2	0	1	18	171	1	0	0	4	1	50
	1	1	2	1	0	0	191	1	0	2	1	25
	1	0	1	1	3	0	0	179	11	0	4	0
	2	0	2	7	0	3	3	4	175	1	3	-25
	1	0	0	3	4	1	1	1	0	186	3	-50
	2	2	2	2	6	1	1	8	1	2	173	-75
	·	·	·	·	·	·	·	·	·	·	·	-100



CA1 - PartB

# Recurrent Neural Network

Problem Statement:

Predicting Movie Review's Sentiment Score

May 20 2025

Kaung Myat San, P2408655

# Agenda

- Background Research & Data Analysis
- Data Augmentation
- Text Preprocessing
- Custom Embedding
- Baseline Model
- Hyper-Tuning
- Conclusion

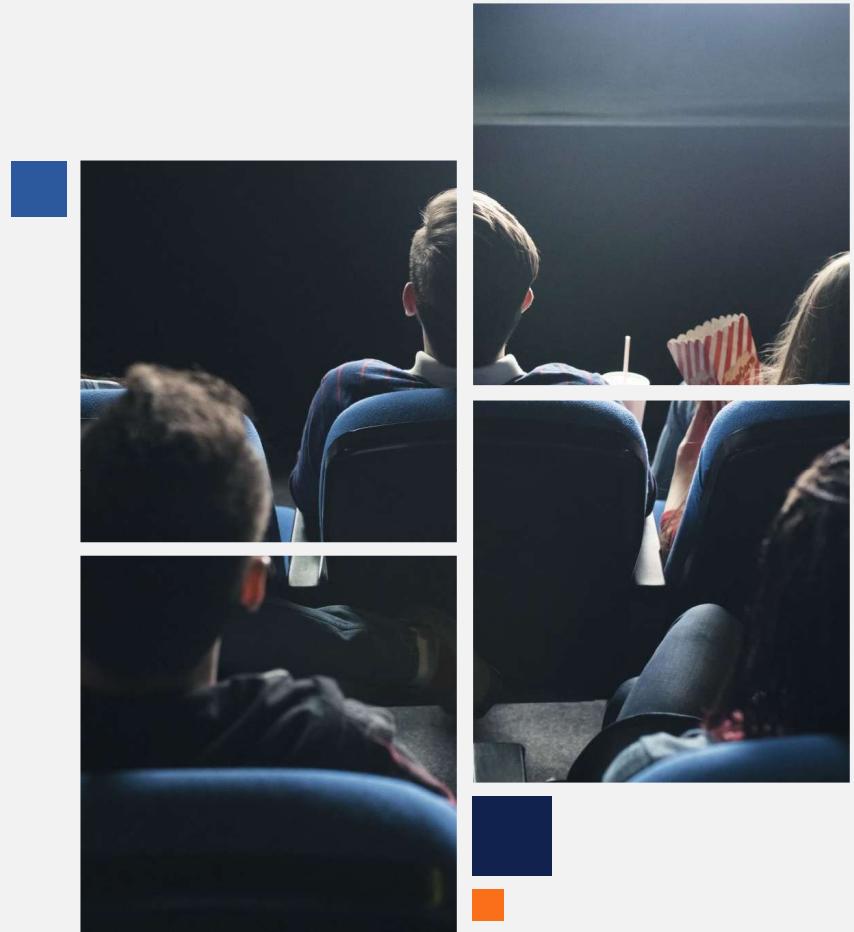




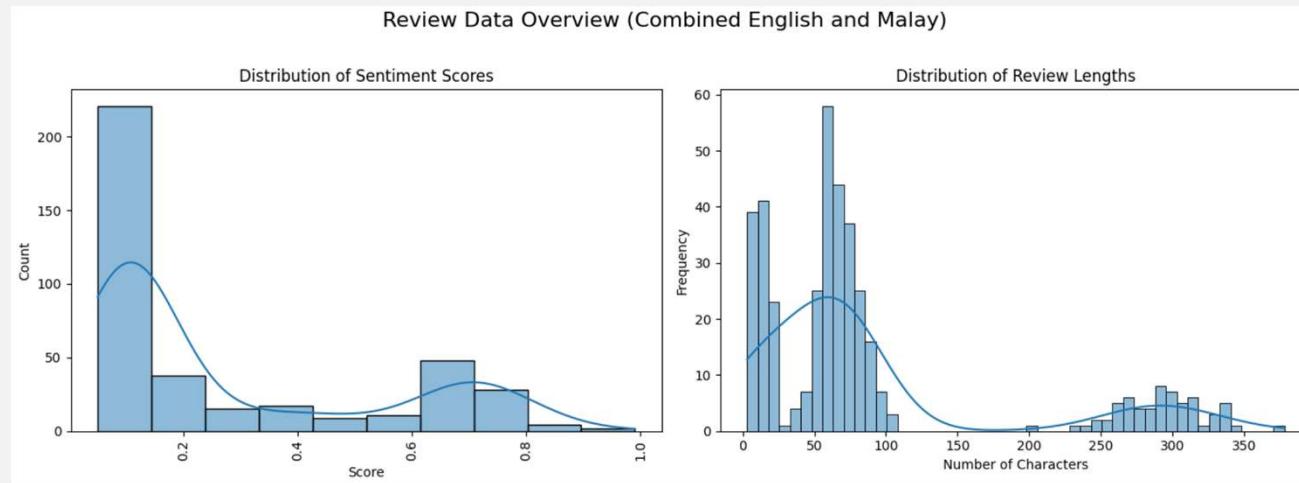
Background  
Research & Data  
Analysis

# Dataset Overview

- Movie Review Dataset contain 4 columns:
  - **Review:** Review of the movie
  - **Score:** Associated Sentiment Score for each Review
  - Are there ways for you to generate more data? Splitting up sentences, would that help?: null
  - **Language:** the language of each review
- **Sentiment Score Range:** 0-1 (positive to negative)
- **No of Language:** 4 (mostly Malay & English)
- **Review:** Convert to English for Normalization and easy interpretation



# Dataset Insights



- Both: Bimodal Distribution, Right-Skewed
- Scores: around 4 times positive as negative
- Review Length: 3 groups – very short, short and long reviews
- Long Reviews: multiple sentences (can split)
- Very Short Reviews: 1, 2 or 4 words (can be used to swap words in Short Reviews)
- Short Reviews: Good for predicting

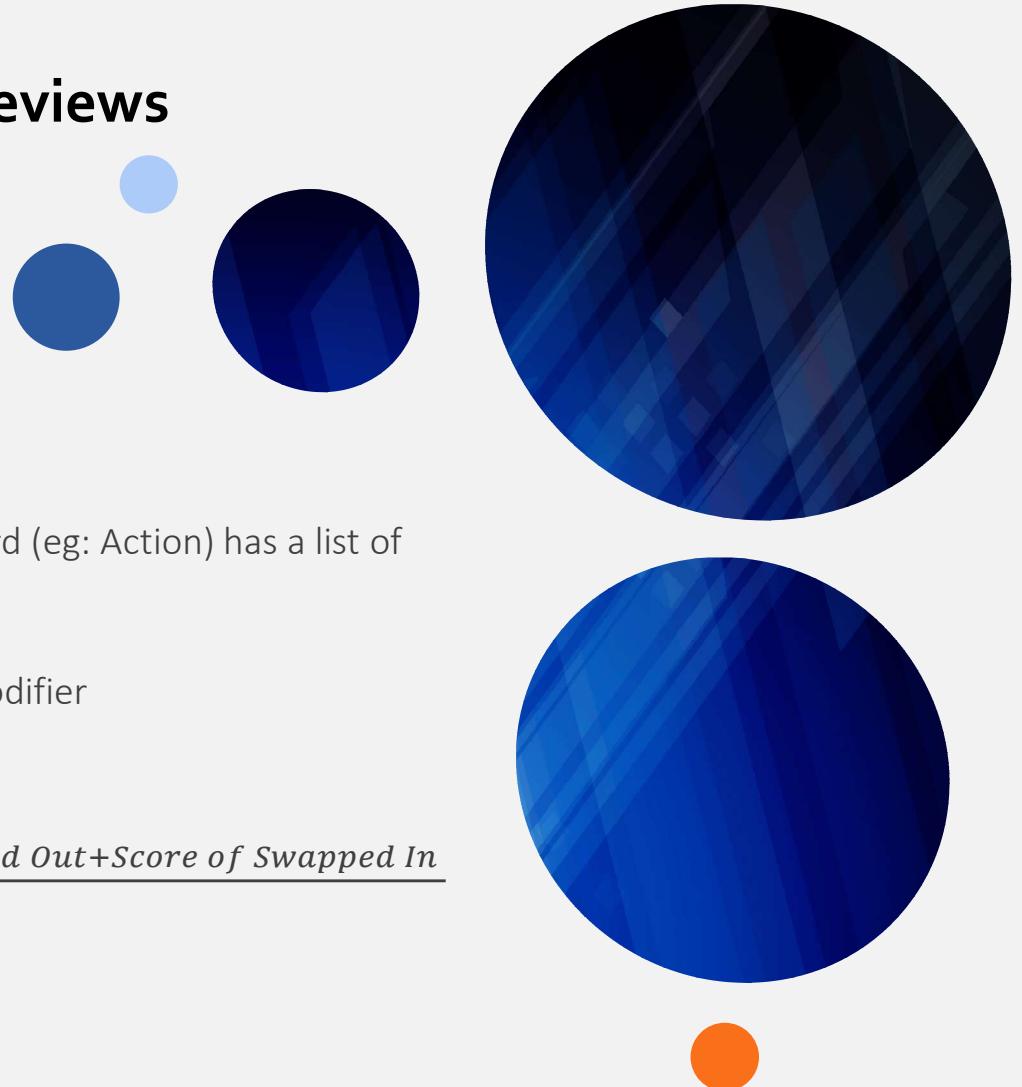




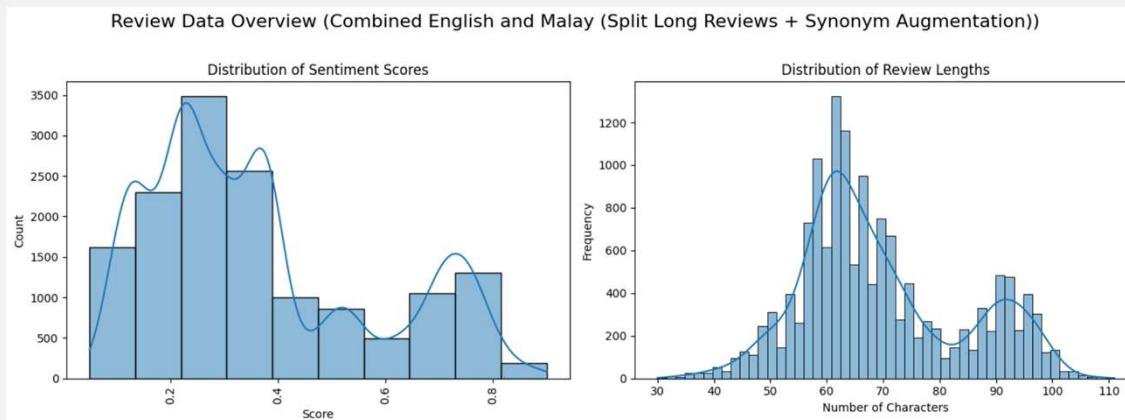
# Data Augmentation

## Reviews Word Swapping Using Reviews (< 30 characters)

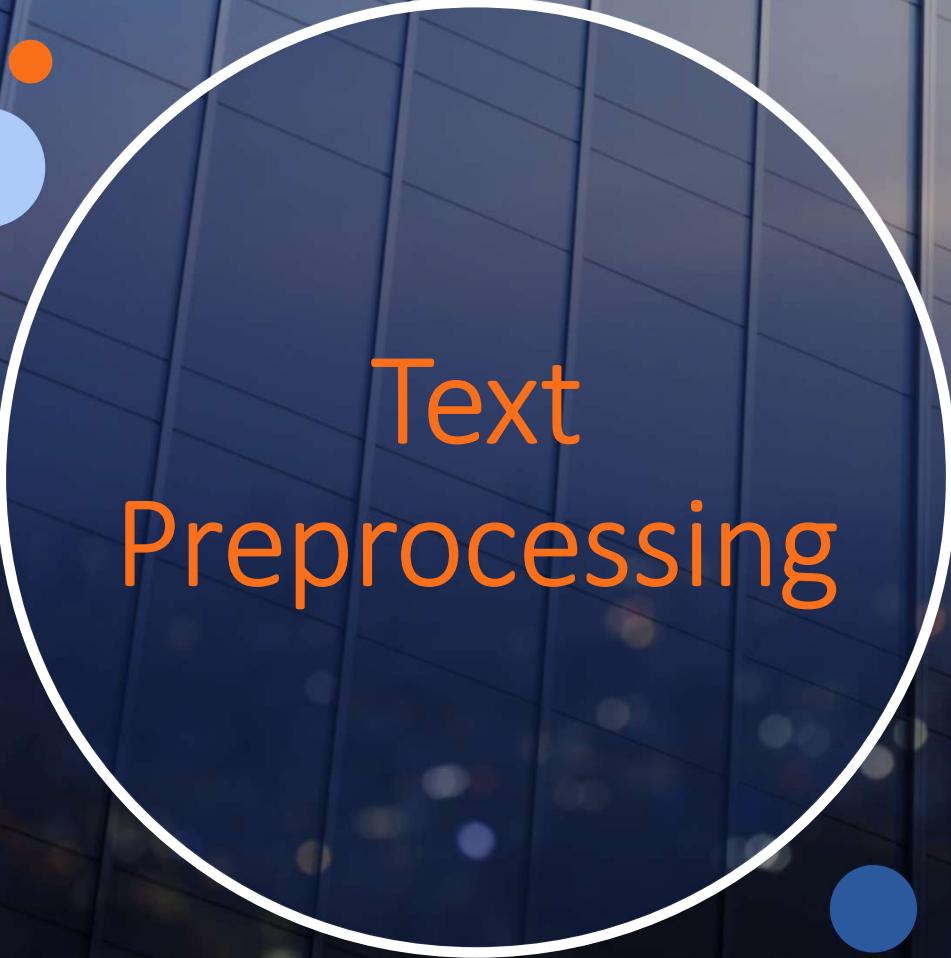
- Reviews (< 30 char) follow a pattern
  - “Amazing”, “Amazing Action”
  - “Breathtaking”, “Breathtaking Visuals”
- Steps to Swap Reviews
  1. Extract Word List with Sentiment Score
  2. Create a Phrase Dictionary where each base word (eg: Action) has a list of modifiers (eg: Amazing)
  3. Find Base Word in Reviews
  4. Swap Modifier if there is a Modifier else Add Modifier
  5. Add this swapped Review as new Rows
  6. Repeat for all Reviews
- New Row Score:  $\frac{\text{Original Score} + \text{Score of Swapped Out} + \text{Score of Swapped In}}{3}$
- Remove Duplicates



# Replace Synonym of Adjective in Reviews



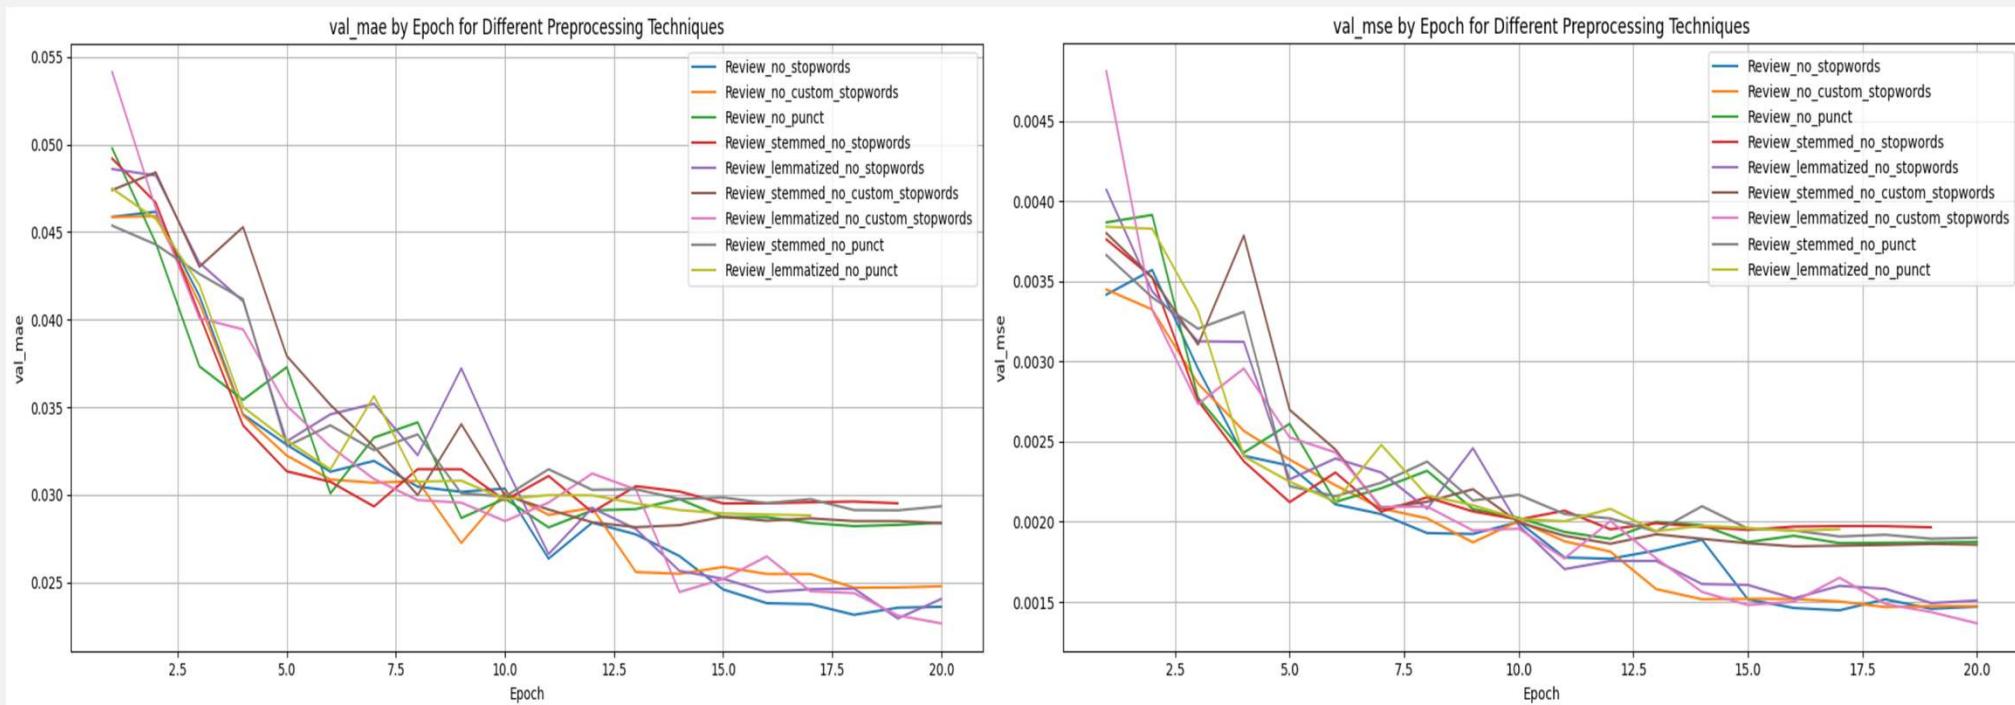
- Create bins and check count in each bins
- To balance dataset, only augment a certain amount in each bin
- Create a new Review with Replaced Synonym
- New Row Score: same score as original
- Remove Duplicates
- Steps to replace Adjective with Synonym
  1. Check if words in Review are adjective
  2. If adjective, find synonym
  3. Return a random synonym
  4. Swap these 2 words
  5. Add this swapped Review as new Rows
  6. Repeat for all Reviews unless amount in bin is reached



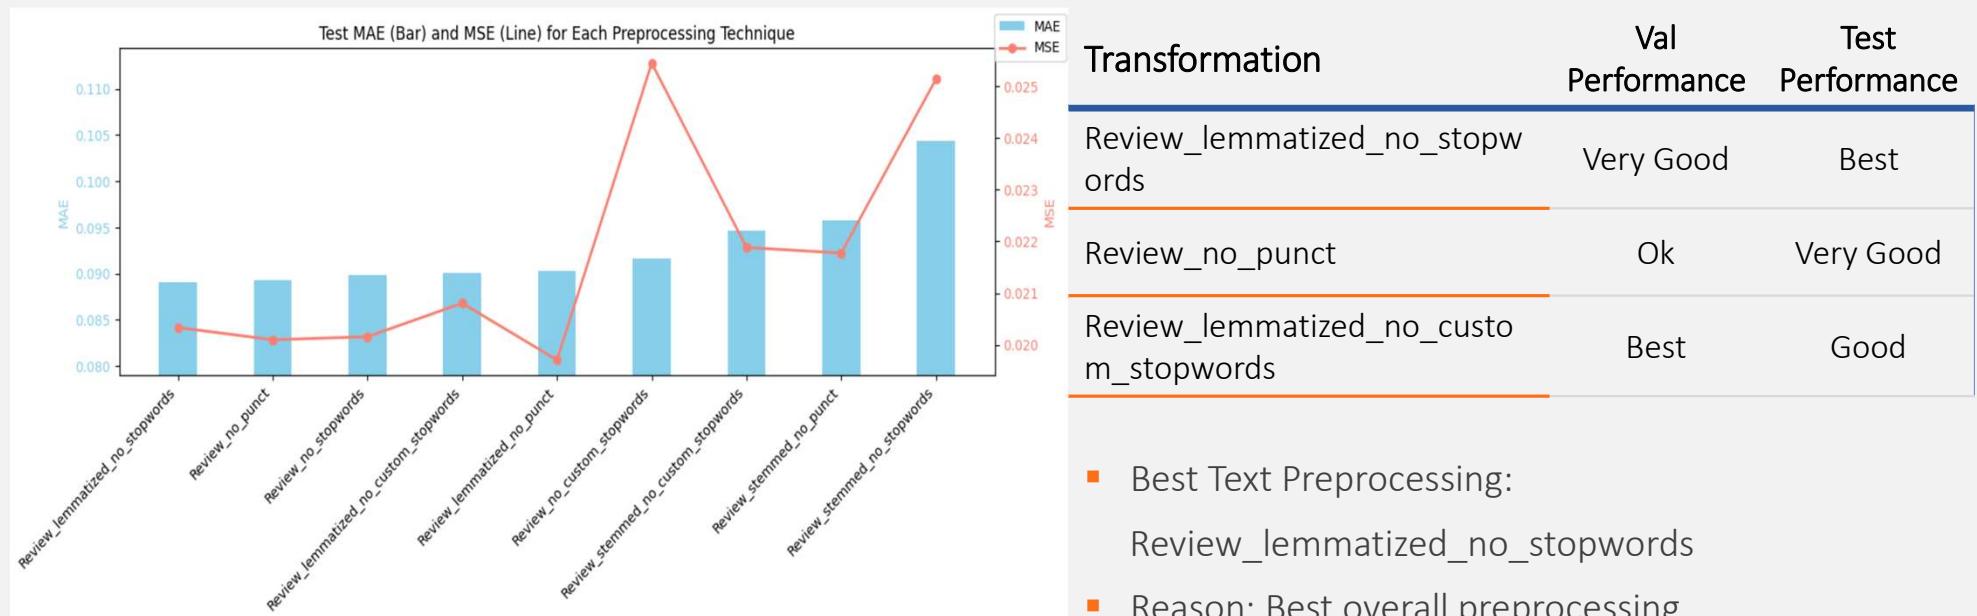
# Text Preprocessing

# Different Text Preprocessing

- All Text Preprocessing have already removed contractions.
- Preprocessing that have removed stop words have also removed punctuation.



# Best Text Transformation

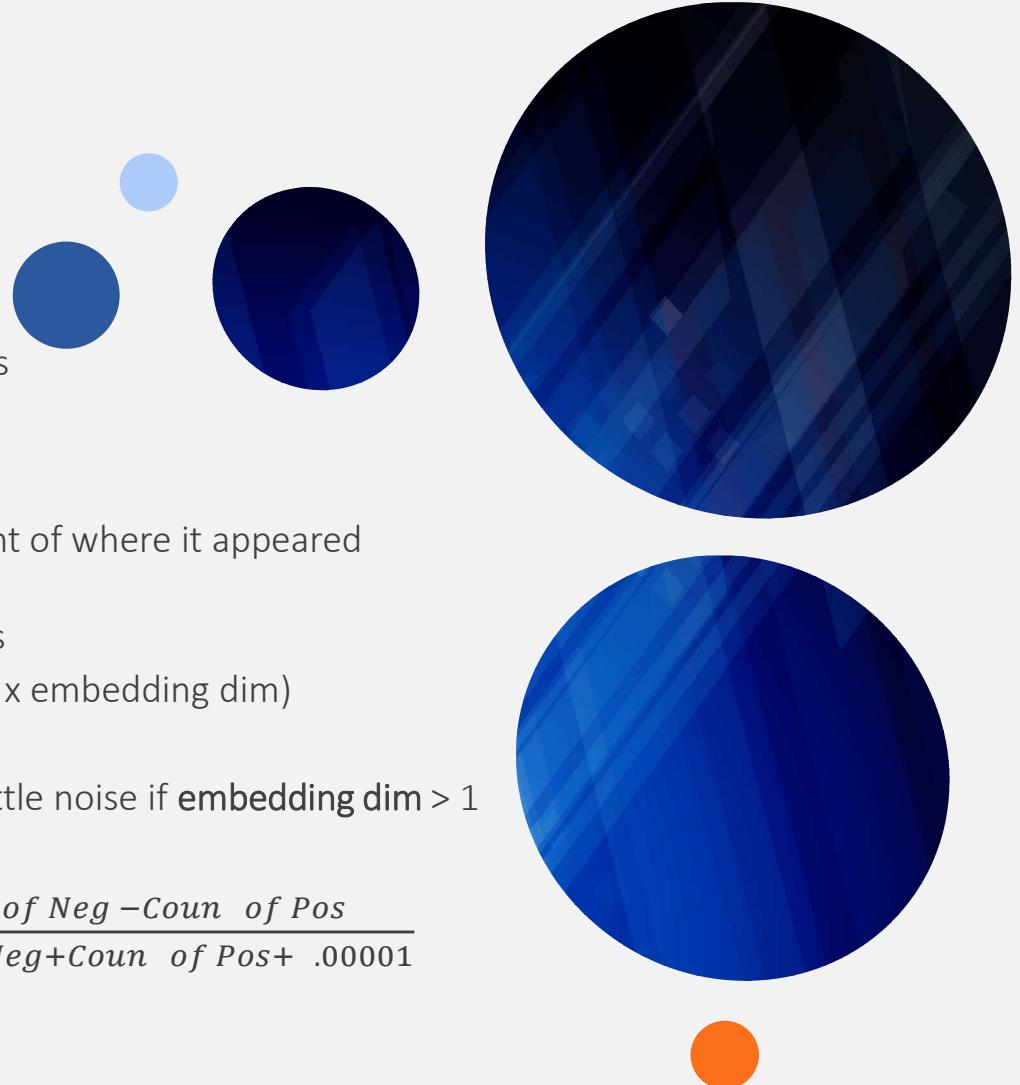




Custom  
Embedding

# Custom Embedding Weights

- Words in Reviews
  - 700 words appeared only in positive reviews
  - 700 words also appeared only in negative reviews
  - Remaining 400 words appeared in both context
- Steps to Create Weights
  1. create a DataFrame where each word has a count of where it appeared positive or negative
  2. tokenize each words to be able to assign weights
  3. Create a embedding matrix (no of unique words x embedding dim)
  4. Calculate embedding weights
  5. Assigned weights to first col and for rest add a little noise if **embedding dim > 1**
  6. Repeat for all words
- Embedding Weights Formula (-1, 1): 
$$\frac{\text{Count of Neg} - \text{Coun of Pos}}{\text{Count of Neg} + \text{Coun of Pos} + .00001}$$





Baseline  
Model

# Baseline Model

- Compile
  - optimizer = Adam(learning\_rate = 0.001)
  - loss = "mse"
- Fitting
  - Early Stopping - monitor= val\_loss, patience = 4, mode = "min"
  - reduce\_lr – monitor = val\_loss, factor = 0.2, patience = 3, mode="min, min\_lr = 1e-6
  - validation\_split = 0.2
  - epochs = 100
  - batch\_size = 32

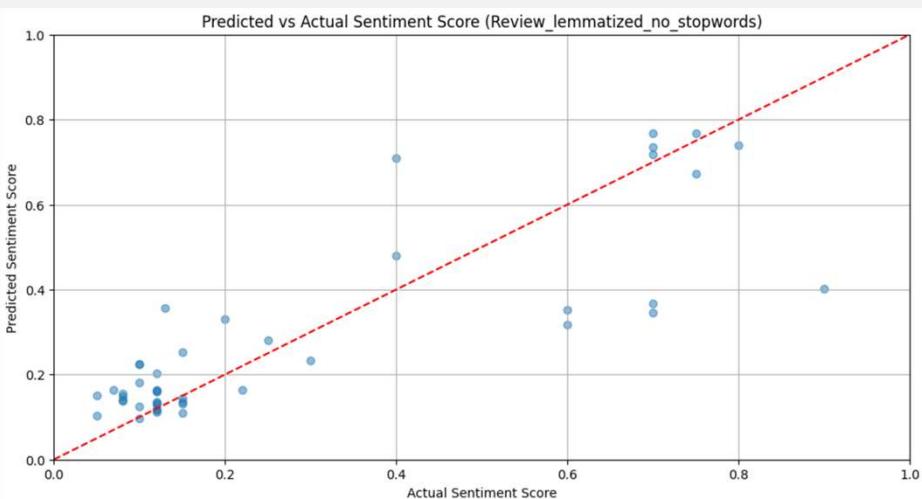
## Test Metrics

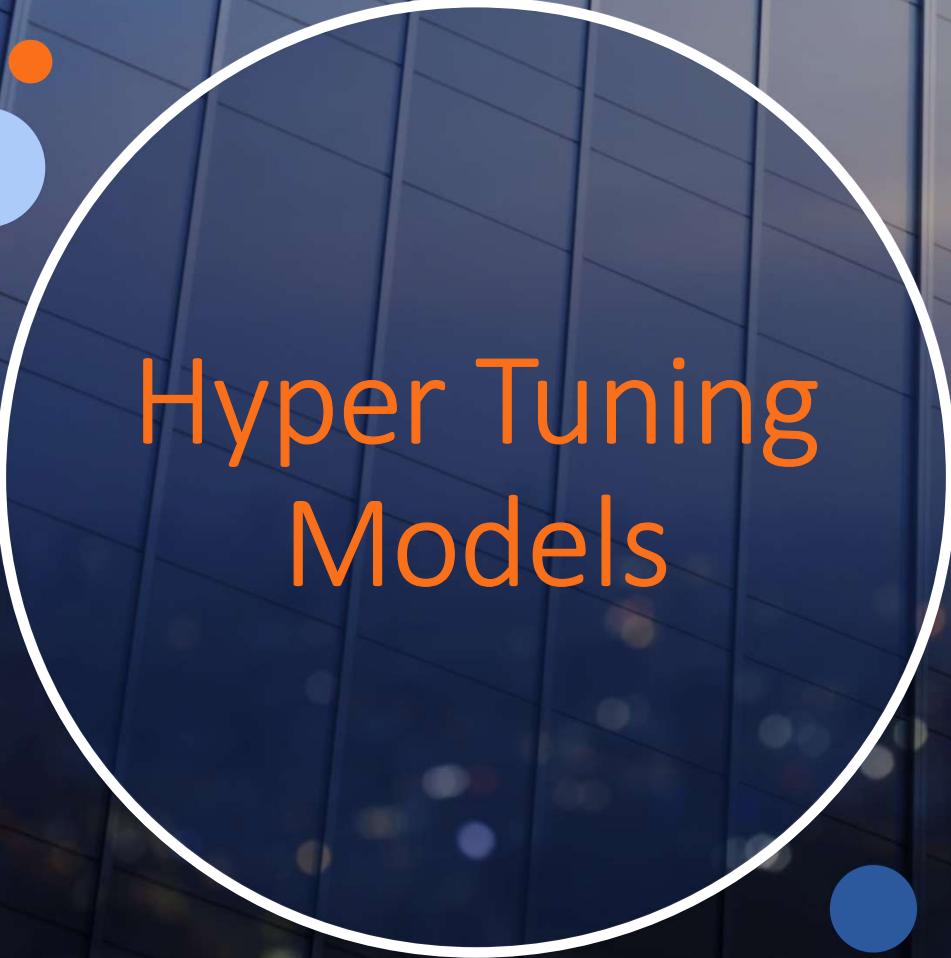
- MAE: 0.0934
- MSE: 0.0206
- MAPE: 47.6619
- MSLE: 0.0097



## Layers                          Parameters

Embedding	weights = custom weight
LSTM	units=64, return_sequence = False
Dense	units=1 , activation="sigmoid"





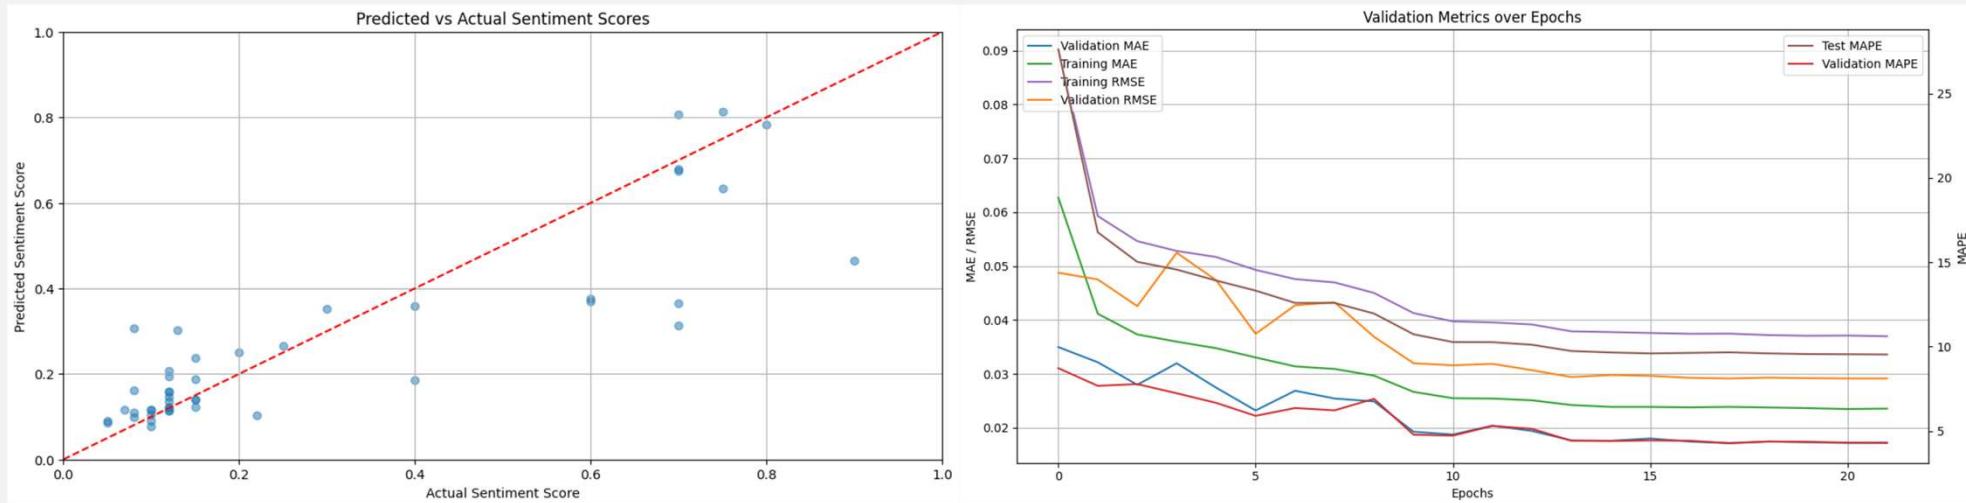
Hyper Tuning  
Models



# Hyper Tuning Parameters

Parameters	Details	Parameters	Details
column_name	name of the column (for vocabulary and embeddings)	dropout_layers	dropout rates for each RNN layer
rnn_type	'LSTM', 'GRU', or 'SimpleRNN'	batch_norm	whether to apply batch normalization
embedding_dim	output dimension of the embedding layer	optimizer	keras optimizer instance
rnn_units	number of RNN units	bidirectional	whether to use bidirectional RNN
No_of_RNN_layers	number of RNN layers	dense_units_list	units for each Dense layers

# Best Hyper Tuned Model

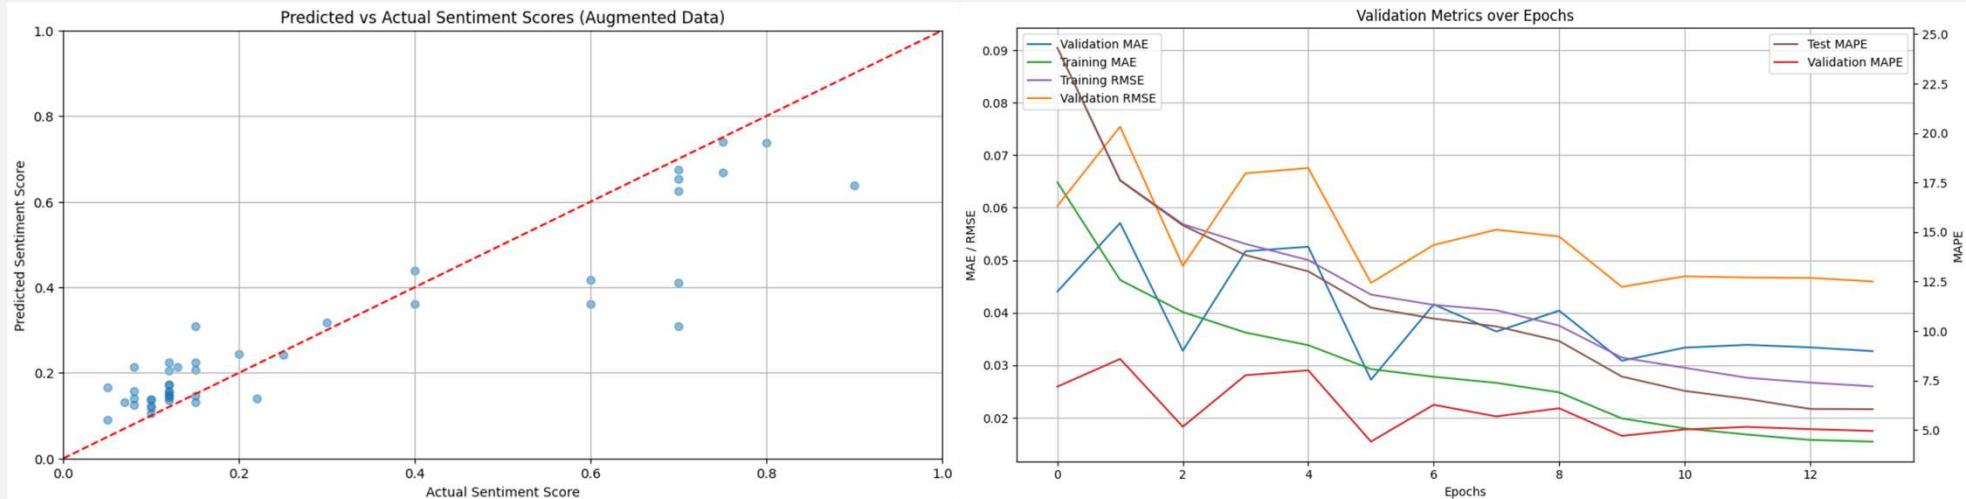


- Test MAE: 0.0808
- Test MSE: 0.0172
- Test RMSE: 0.1311
- Test MSLE: 0.0082

## Model Insights

- Performance:** Great (Low Error Metrics)
- Good Generalization:** Doesn't overfit and underfit
- Prediction:** Perform moderately for Positive Review and ok for Negative Review
- How to Improve:** Augment more Negative Reviews (0.4-1.0)

# Final Best Model



- Test MAE: 0.0758
- Test MSE: 0.0122
- Test RMSE: 0.1106
- Test MSLE: 0.0061

## Model Insights

- Performance: Great (Low Error Metrics)
- Good Generalization: Doesn't overfit and underfit
- Prediction: Perform great for Positive Review and better for Negative Review